

Notes on Inheritance and Subtyping in Java

Prakash Panangaden

1st April 2019

1 Inheritance

One of the fundamental advances in the object-oriented paradigm is the ability to *reuse* code. It often happens that you find yourself coding a small variation of something that you had coded before. If your code is organized into classes, you might observe the following patterns. The new code that you want to write is a new class, but it looks just like an old class that you wrote a while ago except for a couple of methods. It is to handle situations like this that we have the notion of inheritance.

The following code is a basic example.

```
class myInt {

    //Instance variable
    private int n;

    //Constructor
    public myInt(int n){
        this.n = n;
    }

    //Instance methods
    public int getval(){
        return n;
    }

    public void increment(int n){
        this.n += n;
    }

    public myInt add(myInt N){
        return new myInt(this.n + N.getval());
    }
}
```

```

    public void show(){
        System.out.println(n);
    }
}

```

This class just has an integer in each object together with some basic methods. It is not a class we would really write. It is here for illustrative purposes only.

Now imagine that we decide to have “integers” made out of complex numbers. These numbers are called gaussian integers used by the great mathematician Gauss for his work in number theory. He discovered that as an algebraic system, these numbers behaved very much like ordinary integers. We might want to extend our ordinary integers to deal with these gaussian integers. Here is the code that does it. The keyword **extends** in the class declaration tells the system that you want all the code that worked for class **myInt** to be present in class **gaussInt**. We say that **gaussInt** *inherits* from **myInt**. We also say that **myInt** is the *superclass* and that **gaussInt** is the *subclass*.

```

class gaussInt extends myInt {

    //Instance variable
    private int m; //Represents the imaginary part

    /* We do not need the real part that is already present because we
       have inherited all the data and methods of myInt. Thus the
       private int n is also present in every instance of a gaussInt. */

    //Constructor
    public gaussInt(int x, int y){
        super(x); //Special keyword
        this.m = y;
    }

    //Instance methods

    //This method is overridden from the superclass
    public void show(){
        System.out.println(
            "realpart is: " + this.getval() + " imagpart is: " + m);
    }

    public int realpart(){
        return getval();
    }
}

```

```

/*The method getval is defined in the superclass. It is not defined
   here but it is inherited by this class so we can use it. */

public int imagpart(){
    return m;
}

//This is an overloaded method
public gaussInt add(gaussInt z){
    return new gaussInt(z.realpart() + realpart(),
                        z.imagpart() + imagpart());
}

public static void main(String[] args){
    gaussInt kreimhilde = new gaussInt(3,4);
    kreimhilde.show();
    kreimhilde.increment(2);
    kreimhilde.show();
}
} //class gaussInt

```

There are a couple of things to note. In the constructor, we first want to use the constructor for the superclass; this is done with the keyword **super**. The **super** keyword invokes the superclass constructor and then continues with whatever is written next. The picture that you should have in mind is that an instance of **gaussInt** *contains* an instance of **myInt**.

Now you really want to inherit some of the methods unchanged, but some methods need to be modified. Obviously the **show** method of **myInt** is no use for **gaussInt**. In the subclass, you can give a new definition for an old method name. Thus in the above class, we have inherited **getval** and **increment** unchanged but we have modified **show**. This is called *overriding*. From class **gaussInt** you cannot use the **show** method of the superclass as it is well hidden.

There is also a more subtle phenomenon called *overloading*. Look at the **add** method in the two classes. At first sight, this looks like overriding. However the types of the arguments expected by the two definitions of the **add** method are different. The types of the arguments are called the *signature* of the method. Now when we use the same name for a method *but with a different signature* we get two different methods *with the same name*. Both methods are available from the subclass. In this case, from the class **gaussInt**, we can use both **add** methods. How does the system know which one to use? It looks at the types of the actual arguments and decides which one to use. Thus if you want to add an ordinary **myInt** to a **gaussInt**, then the **add** method of the superclass is used.

There are many more subtle aspects to inheritance that you can go into in more detail in later courses. For now, it suffices to grasp the basic concept of inheritance and to recognize

when to use it.

2 Subtyping and Interfaces

One of the features of object-oriented languages is a sophisticated type system that offers a possibility called “subtype polymorphism.” In this section we will explain what this is and how it works in Java. One very important caveat when reading the extant literature: there is terrible confusion about the words “subtyping” and “inheritance”. **They are not the same thing!** Yet one sees in books phrases like “inheritance (i.e. subtyping) is very easy ...”! My only explanation is that lots of people really are clueless. Do not learn from them.

Recall that a type system classifies values into collections that reflect some structural or computational similarity. There is no reason that this classification should be into disjoint collections. Thus, for example, a value like `[]` can be a list of any type. This is the polymorphism that we are familiar with in F#.

The kind of polymorphism that one sees in Java is called “subtype polymorphism” which is based on the idea that there may be a relation between types called *the subtyping relation*. Recall the discussion from the last class. We will say that a type A is a subtype of a type B if *whenever* the context requires an element of type B it can accept an element of type A . We write $A \triangleleft B$ to indicate this.

How do we set up the subtyping relation? In Java, there are some built in instances of subtyping – such as $\text{int} \triangleleft \text{float}$ – but, clearly, this is not worth making such a fuss about. Where subtyping really comes into its own is with user-defined types. In Java, subtyping occurs automatically when you have inheritance; *this does not mean that subtyping and inheritance are the same thing*. You can also have instances of subtyping without any inheritance; this is what happens when you *implement an interface*.

Thus, if we declare a class B to be an extension of class A , we will have - in addition to all the inheritance - that $B \triangleleft A$. In other words, if at some point you have a method `foo` which expects an argument of type A ; `foo` will always accept an object of type B . If we extend class B with class C then we have

$$C \triangleleft B \triangleleft A.$$

If we extend A with another class D then we will have $D \triangleleft A$ but there will be no subtyping relation between B and D or C and D . A method expecting an object of type A will accept objects of type B , C or D . This gives increased generality to the code.

It is often the case that the inheritance hierarchy is not very “wide”. In other words it is unlikely that a class A can be sensibly extended in many incompatible ways. This is because the code has to be inherited and usually only a few methods are modified. If we are modifying all or almost all the methods then it is clear that we are not really using inheritance. We are really trying to get the generality offered by subtype polymorphism.

Java supports subtype polymorphism independently of inheritance. This is done by **interfaces**. An interface is a declaration of a collection of method names - *without any method bodies* - and perhaps some constants. They describe a (sub)set of methods that

a class might have. There is no code in an interface, hence there is nothing to inherit. A concrete class¹ may be declared to **implement** an interface. This is really a subtyping declaration. An interface names a new type (not a class) and when we say that a class P implements interface I we have set up the subtyping relation $P \triangleleft I$. Because there was no code in the interface, P does not inherit any code; but *it must have a method of the same name and type as every method name in the interface I .*

A class can implement many interfaces and thus one can have a complex type hierarchy with multiple subtyping. One often hears the phrase “interfaces allow Java to fake multiple inheritance”. This is a source of confusion. What interfaces allow is multiple subtyping. Java definitely does not have multiple inheritance (C++ does have true multiple inheritance); what it has is multiple subtyping.

Here is an example of the use of inheritance. Imagine that you have written code to draw the graph of a mathematical function. You want this to be abstracted on the function. You do not want to write code to plot a graph of the *sine* function and another different - but almost identical - piece of code to plot a graph of the *exp* function. You want the function - a piece of code - to be a *parameter*. We can do this with objects because objects can be passed around like any other piece of data but yet they carry code. What kind of object is a mathematical function? It expects a **double** argument and returns a **double** result. There might be other methods associated with function objects but the **plot** method does not care. It only cares that there is a method - called, say, **y** - such that **f.y(<input>)** returns a double. So instead of having to know all about the details of the class of mathematical functions we just define an interface **plottable** with the one method **y** in it with the appropriate type. When we define our mathematical function objects we can make them as complicated as we please as long as we declare that they implement **plottable** and ensure that they really do have the method **y**. If we have another type of object - say **fin-data** - for financial data we would expect to define a totally different class with no relation to mathematical function objects. However, **fin-data** could also have a method **y** and be declared to implement **plottable**. Then our **plot** method works on totally unrelated classes. We have achieved generality for our plotting method through subtype polymorphism in a very general situation. Of course, we could have done this with inheritance too but it would be silly.

In between interfaces and classes are *abstract classes* that have some methods defined and some that are left blank as in interfaces. You can extend them as you would any class.

There are many subtle issues in Java about method lookup and how it interacts with inheritance and type-checking. We will deal with them in the next section.

3 Subtypes and Typechecking in Java

In this section we revisit the gaussian integers example and formalize the rules for type checking a little more carefully.

The following is a brief summary of the

¹We use the adjective “concrete” to mean that the class is completely defined, i.e. it has all the actual code for the methods.

Rules for Method Lookup and Type Checking.

First the rules. There are two phases: **compile time**, which is when type checking is done and **run time**, which is when method lookup happens. Compile time is *before* run time.

- The type checker has to say that a method call is OK at compile time.
- All type checking is done based on what the declared type of a reference to an object is.
- Subtyping is an integral part of type checking. This means if B is a subtype of A and there is a context that gets a B where A was expected there will not be a type error.
- Method lookup is based on actual type of the object and not the declared type of the reference.
- When there is overloading (as opposed to overriding) this is resolved by type-checking.

3.1 Analysis of the gaussInt Example

Here is an example exhibiting both overriding and overloading.

```
class myInt {
    private int n;
    public myInt(int n){ this.n = n;}

    public int getval(){return n;}

    public void increment(int n){ this.n += n;}

    public myInt add(myInt N){ return new myInt(this.n + N.getval());}

    public void show(){
        System.out.println(n);}
}

class gaussInt extends myInt {
    private int m; //represents the imaginary part

    public gaussInt(int x, int y){
        super(x);
        this.m = y;}
}
```

```

public void show(){
    System.out.println(
        "realpart is: " + this.getval() + " imagpart is: " + m);}

public int realpart() {return getval();}
public int imagpart() {return m;}

public gaussInt add(gaussInt z){
    return new gaussInt(z.realpart() + realpart(),
                        z.imagpart() + imagpart());
}

public static void main(String[] args){
    gaussInt k = new gaussInt(3,4);
    k.show();
    k.increment(2);
    k.show();

    System.out.println("Now we watch the subtleties of overloading.");
    myInt a = new myInt(3);
    gaussInt z = new gaussInt(3,4);
    gaussInt w;    //no object has been created
    myInt b = z;    //b and z are names for the same object
    //even though b and z refer to the SAME object they have
    //different types

    System.out.print("the value of z is: "); z.show();
    System.out.print("the value of b is: "); b.show();
    //which show method will be used?

    myInt d = b.add(b); //this does type check
    System.out.print("the value of d is: "); d.show();
    // w = z.add(b); will not type check
    // w = b.add(z); will not type check
    w = ( gaussInt) b.add(z); //this does type check
    System.out.print("the value of w is: "); w.show();
    myInt c = z.add(a); //will this typecheck?
    System.out.print("the value of c is: "); c.show();

}
}

```

Here is what we can say about the variables (TBD means “to be determined”):

Name	Declared Type	ActualType
a:	myInt	myInt
z	gaussInt	gaussInt
w	gaussInt	TBD
b	myInt	gaussInt
d	myInt	myInt
c	myInt	TBD

```
myInt a = new myInt(3);
gaussInt z = new gaussInt(3,4);
gaussInt w;
myInt b = z;
```

```
System.out.println("the value of z is"+ z.show());
```

```
> real part is 3 imag part is 4
```

this prints out the above line because *z* is declared to be of type `gaussInt`. It passes the type checker as there is a `show` method defined in the `gaussInt` class. At run time it uses the `show` method of `gaussInt` to display the above line.

```
System.out.println("the value of b is :" + b.show());
```

```
> real part is 3 and imag part is 4.
```

b is declared to be of type `myInt`. There is a method called `show` in the `myInt` class. The type checker sees that and because of that it passes the type checker, **but** the actual type of *b* is `gaussInt`. Method lookup is based on actual types of objects and therefor *b* uses the `show` method in the `gaussInt` class and displays what a `gaussInt` object would have shown.

```
myInt d = b.add(b)
```

```
System.out.println("the value of d is:"+ d.show());
```

```
> 6
```


b is declared to be of the type `myInt`, the type checker checks to see whether there is an `add` method in the `myInt` class. Yes there is one; it takes a `myInt` object and returns a `myInt` object as the result. At run time b 's actual type is `gaussInt` the run-time system checks to see if there is an `add` method in the `gaussInt` class which matches the type that it was told by the type-checker. There are two `add` methods - one that takes a `myInt` and returns a `myInt` (This method has been inherited from the `myInt` class). The other takes a `gaussInt` and returns a `gaussInt`; this is the method that is explicitly defined in the `gaussInt` class. However the latter method does not match what the type-checker told the run-time system to expect.

NOW WHICH ADD METHOD DO WE USE?

since "When there is overloading, it is resolved by typechecking" the method which takes an object of the type `myInt` will be used. This is the method that has been inherited. It takes in a `myInt` and returns a `myInt`. Hence `b.add(b)` returns a `myInt` object and therefor NOW the actual type of d is `myInt`.

```
//w= z.add(b) -----(i)
//w = b.add(z)----- (ii)
```

These two will not type check

1. z is declared to be of the type `gaussInt`. There are two methods in the `gaussInt` class, the one that takes in a `myInt` object and returns a `myInt` object is used. Why? Once again overloading is resolved by typechecking. Since b is declared to be a `myInt` object it will pick the `add` method that it inherited.

z is a `gaussInt` which is a subtype of `myInt` and hence is added to b and returns a `myInt`. w is declared to be a `guassInt`. Since `myInt` is not a subtype of `gaussInt` the assignment statement will not accept this for the right hand side, and hence would cause an error.

2. b is declared to be of the type `myInt`. The type checker checks if there is an `add` method in the `myInt` class there is one which expects a `myInt` object and returns a `myInt` object. z is a `gaussInt` and since `gaussInt` is a subtype of `myInt`, b is added to z to produce a `myInt` object

w is declared to be a `gaussInt` Since `myInt` is not a subtype of `gaussInt` it will not accept it and hence would cause an error.

```
w = ( (gaussInt) b).add(z)
```

This does type check as it is just a little modification to case 2 above. Now since w is a `gaussInt`, it better get a `gaussInt` on the right hand side. However, now, because of the cast, the type-checker knows that b is really a `gaussInt`. Thus, it now has to choose between two possible `add` methods. To resolve the overloading it uses the declared types; z has declared type `gaussInt`. Thus when it resolves the overloading of the `add` method it figures out to use the `gaussInt` to `gaussInt` version.